

NLP Assignment-4

Name:P.Rahul

Roll No:2203A51439

Batch:12

1. Load data from keras.datasets and perform following computational analysis:- [CO2]
 - (a) Preprocessing of the Data
 - (b) Divide data into training and testing data set
 - (c) Build the Recurrent Neural network (RNN) Model
 - (d) Training the RNN Model
 - (e) Evaluate the model on the test dataset to see how well it generalizes.
2. Develop LSTM (Long Short-Term Memory) by utilizing data set from <https://www.kaggle.com/code/amirrezaeian/time-series-data-analysis-using-lstm-tutorial> Links to an external site. or take any time series data. [CO2]
3. Demonstrate Vanishing and Exploding Gradients on deep neural network. [CO2]

Task-1:

(a) Preprocessing of the Data

```
Task 1: RNN with IMDB Dataset

from keras.datasets import imdb
from keras.preprocessing import sequence

# Load the dataset
max_words = 10000 # Top 10,000 most frequent words in the dataset
maxlen = 500 # Limit the review length to 500 words

(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_words)

# Preprocessing: Pad sequences to ensure uniform length
x_train = sequence.pad_sequences(x_train, maxlen=maxlen)
x_test = sequence.pad_sequences(x_test, maxlen=maxlen)

print(f'Training data shape: {x_train.shape}')
print(f'Testing data shape: {x_test.shape}')
```

Training data shape: (25000, 500)
Testing data shape: (2500, 500)

(b) Divide data into training and testing data set

The data is already divided into training and testing data set

(c) Build the Recurrent Neural network (RNN) Model

```
from keras.models import Sequential
from keras.layers import SimpleRNN, Dense, Embedding

# RNN Model
model = Sequential()
model.add(Embedding(max_words, 128)) # Embedding layer to convert words into vectors
model.add(SimpleRNN(64)) # Simple RNN with 64 units
model.add(Dense(1, activation='sigmoid')) # Output layer with a single neuron for binary classification

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	?	0 (unbuilt)
simple_rnn_1 (SimpleRNN)	?	0 (unbuilt)
dense_1 (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)
Trainable params: 0 (0.00 B)
Non-trainable params: 0 (0.00 B)

(d) Training the RNN Model

```
# Train the RNN
history = model.fit(x_train, y_train, epochs=5, batch_size=64, validation_split=0.2)
```

Epoch 1/5
313/313 — 78s 241ms/step - accuracy: 0.6057 - loss: 0.6482 - val_accuracy: 0.7004 - val_loss: 0.5710
Epoch 2/5
313/313 — 79s 252ms/step - accuracy: 0.7617 - loss: 0.5027 - val_accuracy: 0.7056 - val_loss: 0.5680
Epoch 3/5
313/313 — 76s 242ms/step - accuracy: 0.8142 - loss: 0.4144 - val_accuracy: 0.8098 - val_loss: 0.4323
Epoch 4/5
313/313 — 75s 241ms/step - accuracy: 0.8895 - loss: 0.2748 - val_accuracy: 0.5812 - val_loss: 0.6646
Epoch 5/5
313/313 — 73s 234ms/step - accuracy: 0.7709 - loss: 0.4844 - val_accuracy: 0.7476 - val_loss: 0.5522

(e) Evaluate the model on the test dataset to see how well it generalizes.

```
[8] # Evaluate the model
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f'Test Accuracy: {test_acc:.4f}')
```

782/782 32s 40ms/step - accuracy: 0.7466 - loss: 0.5580
Test Accuracy: 0.7495

Task-2:

Task 2: Long Short-Term Memory (LSTM) for Time Series Data

```
from keras.models import Sequential
from keras.layers import LSTM, Dense
import numpy as np

# Simulating a time series dataset (random for demonstration)
n_timesteps = 100
n_features = 1
X = np.random.randn(1000, n_timesteps, n_features)
y = np.random.randn(1000, 1)

# Split into training and testing sets
split = int(0.8 * len(X))
X_train, X_test = X[:split], X[split:]
y_train, y_test = y[:split], y[split:]

# LSTM Model
model = Sequential()
model.add(LSTM(64, input_shape=(n_timesteps, n_features)))
model.add(Dense(1))

model.compile(optimizer='adam', loss='mse')
model.summary()
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/rnn.py:204: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an
super().__init__(**kwargs)
Model: "sequential_2"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 64)	16,000
dense_2 (Dense)	(None, 1)	65

Total params: 16,065 (66.25 KB)
Trainable params: 16,065 (66.25 KB)
Non-trainable params: 0 (0.00 B)

```
[12] # Train the LSTM Model
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=0.2)
```

Epoch 1/10
20/20 3s 54ms/step - loss: 0.9358 - val_loss: 1.0835
Epoch 2/10
20/20 1s 40ms/step - loss: 0.8912 - val_loss: 1.0816
Epoch 3/10
20/20 1s 42ms/step - loss: 0.9875 - val_loss: 1.0838
Epoch 4/10
20/20 1s 41ms/step - loss: 0.9475 - val_loss: 1.0744
Epoch 5/10
20/20 1s 41ms/step - loss: 0.8904 - val_loss: 1.0842
Epoch 6/10
20/20 1s 42ms/step - loss: 0.8957 - val_loss: 1.0789
Epoch 7/10
20/20 1s 64ms/step - loss: 0.9805 - val_loss: 1.0757
Epoch 8/10
20/20 1s 67ms/step - loss: 0.9886 - val_loss: 1.0868
Epoch 9/10
20/20 1s 67ms/step - loss: 0.9439 - val_loss: 1.0811
Epoch 10/10
20/20 1s 51ms/step - loss: 0.9527 - val_loss: 1.0817

```
[13] # Evaluate the model
test_loss = model.evaluate(X_test, y_test)
print(f'Test Loss: {test_loss:.4f}')
```

7/7 0s 16ms/step - loss: 0.9977
Test Loss: 0.9880

Task-3:

Task 3: Demonstrate Vanishing and Exploding Gradients in Deep Neural Networks

```
from keras.models import Sequential
from keras.layers import Dense
from keras.optimizers import SGD
import tensorflow as tf

# Build a deep neural network with many layers
model = Sequential()
model.add(Dense(64, activation='tanh', input_shape=(n_timesteps,)))

# Add many layers to show the vanishing/exploding gradient problem
for _ in range(50): # 50 layers
    model.add(Dense(64, activation='tanh'))

model.add(Dense(1))

# Using Stochastic Gradient Descent (SGD) to highlight the gradient problem
optimizer = SGD(learning_rate=0.01)
model.compile(optimizer=optimizer, loss='mse')
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a `Layer`.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```
[17] import tensorflow as tf

# Custom Callback to Monitor Gradients
class GradientMonitor(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs=None):
        # Manually select a batch from the training data for gradient monitoring
        batch_data = X_train[:32] # Take the first batch (you can randomize this)
        batch_labels = y_train[:32]

        with tf.GradientTape() as tape:
            # Forward pass
            predictions = self.model(batch_data, training=True)
            loss_value = self.model.compiled_loss(batch_labels, predictions)

            # Compute gradients
            gradients = tape.gradient(loss_value, self.model.trainable_weights)

            # Extract max and min gradients
            max_grad = max([tf.reduce_max(tf.abs(g)) for g in gradients if g is not None])
            min_grad = min([tf.reduce_min(tf.abs(g)) for g in gradients if g is not None])

            print(f'Epoch {epoch+1}: Max Gradient = {max_grad:.6f}, Min Gradient = {min_grad:.6f}')

# Train the model and monitor gradients
history = model.fit(X_train, y_train, epochs=5, batch_size=32, validation_split=0.2, callbacks=[GradientMonitor()])
```

Epoch 1/5
11/20 — 0s 5ms/step - loss: 0.8358 Epoch 1: Max Gradient = 0.336394, Min Gradient = 0.000000
20/20 — 0s 21ms/step - loss: 0.8820 - val_loss: 1.0989
Epoch 2/5
/usr/local/lib/python3.10/dist-packages/keras/src/backend/tensorflow/trainer.py:607: UserWarning: `model.compiled_loss()` is deprecated. Instead, use `model.compute_loss(x, y, y_pred, sample_weight, training)`.
warnings.warn(
20/20 — 0s 5ms/step - loss: 0.9715 Epoch 2: Max Gradient = 0.436690, Min Gradient = 0.000000
20/20 — 1s 17ms/step - loss: 0.9703 - val_loss: 1.1001
Epoch 3/5
11/20 — 0s 5ms/step - loss: 0.8555 Epoch 3: Max Gradient = 0.447012, Min Gradient = 0.000000
20/20 — 0s 15ms/step - loss: 0.8866 - val_loss: 1.1174
Epoch 4/5
11/20 — 0s 6ms/step - loss: 0.9292 Epoch 4: Max Gradient = 0.211720, Min Gradient = 0.000000
20/20 — 1s 28ms/step - loss: 0.9288 - val_loss: 1.1172
Epoch 5/5
11/20 — 0s 12ms/step - loss: 0.8833 Epoch 5: Max Gradient = 0.390337, Min Gradient = 0.000000
20/20 — 1s 35ms/step - loss: 0.8835 - val_loss: 1.1101